

Deep Learning for Computer Vision

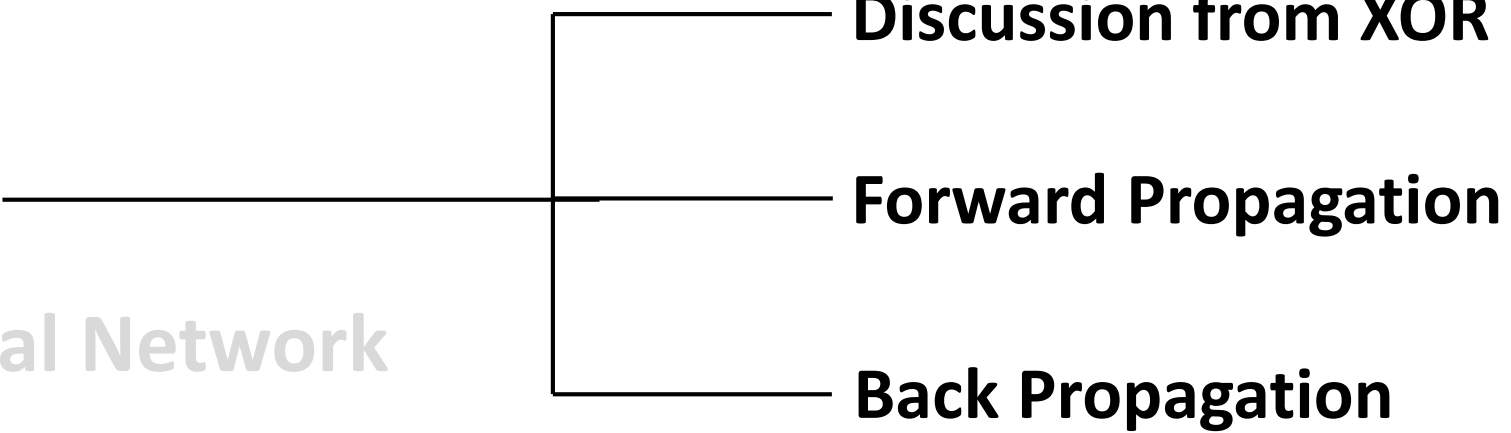
Libao Zhang, Jie Ma

2017.11.28

Content

- **Neural Network**
- **Convolution Neural Network**
- **Deep Learning for Computer Vision**

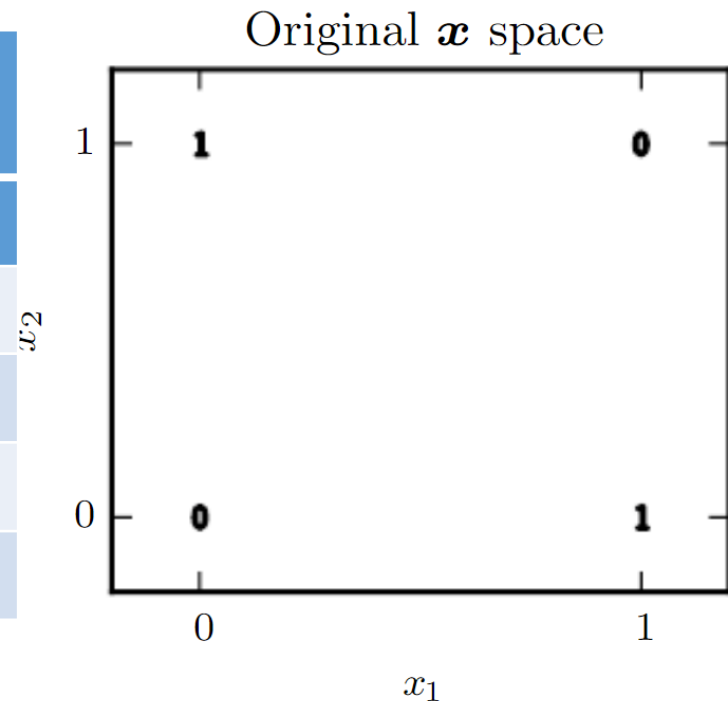
Content

- **Neural Network** ———— 
 - Discussion from XOR**
 - Forward Propagation**
 - Back Propagation**
- Convolution Neural Network
- Deep Learning for Computer Vision

Neural Network: Discussion from XOR

- The **XOR function** (“exclusive or”) is an operation on two **binary values**, x_1 and x_2 . When **exactly one** of these binary values is equal to 1, the XOR function returns 1. Otherwise, it returns 0.

Input		Target Output
x_1	x_2	y
0	0	0
1	1	0
1	0	1
0	1	1



Neural Network: Discussion from XOR

- We can treat this problem as a **regression problem** and use a **mean squared error(MSE)** loss function.
- Target function:

$$y = f^*(x),$$

where $f^*(1, 0) = 1, f^*(0, 1) = 1, f^*(1, 1) = 0, f^*(0, 0) = 0$

- Our model:

$$y = f(x, \vartheta)$$


- MES loss function:

$$J(\theta) = \frac{1}{4} \left(\sum_{i=1}^4 [f(x_i, \vartheta) - f^*(x_i)]^2 \right)$$

Neural Network: Discussion from XOR

- Suppose that we choose a **linear model**, with θ consisting of w and b . Our model is defined to be:

$$\begin{aligned}f(x; w, b) &= x_1 * w_1 + x_2 * w_2 + b \\&= (w_1, w_2) * \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + b \\&= Wx + b\end{aligned}$$




- Minimize $J(\theta)$, we obtain $W = 0, b = \frac{1}{2}$

Input		Target Output	Regression by linear model
x_1	x_2	y^*	y
0	0	0	0.5
1	1	0	0.5
1	0	1	0.5
0	1	1	0.5

Neural Network: Discussion from XOR

- We must use a **nonlinear function** to describe the features.




- The network now contains two function:


$$\begin{aligned} \begin{bmatrix} h1 \\ h2 \end{bmatrix} &= [x_1 * w_{11} + x_2 * w_{12}, x_1 * w_{21} + x_2 * w_{22}]^T + b \\ &= \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{bmatrix} * \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + b = wx + b \end{aligned}$$

$$y = h1 * W_1 + h2 * W_2 + c = Wh + c$$

Neural Network: Discussion from XOR



Rectified linear unit
 $g(z) = \max\{0, z\}$



- The whole model:

$$f(x; W, c, w, b) = W * \max\{0, w * x + b\} + c$$

- We can then specify a solution to the XOR problem:

$$w = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} \quad W = \begin{bmatrix} 1 \\ -2 \end{bmatrix}$$

$$b = \begin{bmatrix} 0 \\ -1 \end{bmatrix} \quad c = 0$$


Input		Target Output	Regression by linear model
x_1	x_2	y^*	y
0	0	0	0
1	1	0	0
1	0	1	1
0	1	1	1

Neural Network: Feedforward Propagation

- There are n_l neurons in layer l , the output of layer l is:

$$\mathbf{a}^{(l)} = [\mathbf{a}_1^{(l)}, \mathbf{a}_2^{(l)}, \dots, \mathbf{a}_{n_l}^{(l)}]^T$$

- Denote the weight associated with the connection between unit j in layer l , and unit i in layer $l + 1$: $\mathbf{W}^{(l)} = [\mathbf{w}_{ij}^{(l)}]_{n_{l+1} \times n_l}$
- The bias associated with unit i in layer $l + 1$: $\mathbf{b}_i^{(l)}$



Neural Network: Feedforward propagation

- **Feedforward propagation algorithm:**

$$z^{(l+1)} = W^{(l)} a^{(l)} + b^{(l)}$$
$$a^{(l+1)} = f(z^{(l+1)})$$

- **Example:**

$$\text{Input: } \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$$

$$\text{Layer1: } W^{(1)} = \begin{bmatrix} 0.1 & 0.7 \\ 0.2 & 0.8 \end{bmatrix}, b^{(1)} = \begin{bmatrix} -1 \\ -0.8 \end{bmatrix}$$

$$\text{Layer2: } W^{(2)} = [0.2 \quad -0.3], b^{(2)} = -0.5$$


- **Feedforward propagation:**

$$z^{(2)} = W^{(1)} a^{(1)} + b^{(1)} = \begin{bmatrix} 0.1 & 0.7 \\ 0.2 & 0.8 \end{bmatrix} * \begin{bmatrix} 1 \\ 2 \end{bmatrix} + \begin{bmatrix} -1 \\ -0.8 \end{bmatrix} = \begin{bmatrix} 0.5 \\ 1 \end{bmatrix}$$

$$a^{(2)} = f(z^{(2)}) = \begin{bmatrix} 0.5 \\ 1 \end{bmatrix}$$

$$z^{(3)} = W^{(2)} a^{(2)} + b^{(2)} = [0.2 \quad -0.3] * \begin{bmatrix} 0.5 \\ 1 \end{bmatrix} + (-0.5) = -0.7$$


$$a^{(3)} = f(z^{(3)}) = 0$$



Neural Network: Back Propagation

- Question:

How to obtain the **appropriate** parameters?



Neural Network: Feedforward propagation

- For a regression or classification task, we have dataset with **input feature** and **ground truth**.
- Divide the dataset into 2 parts, one for **training** and one for **testing**.
- In **training set** D , we want to minimize the difference between **the label of training data and the model's predictions**.

- MSE loss function:

$$J(\theta) = \sum_{i \in D} \frac{1}{2} \|f(x_i; \theta) - y_i\|^2$$

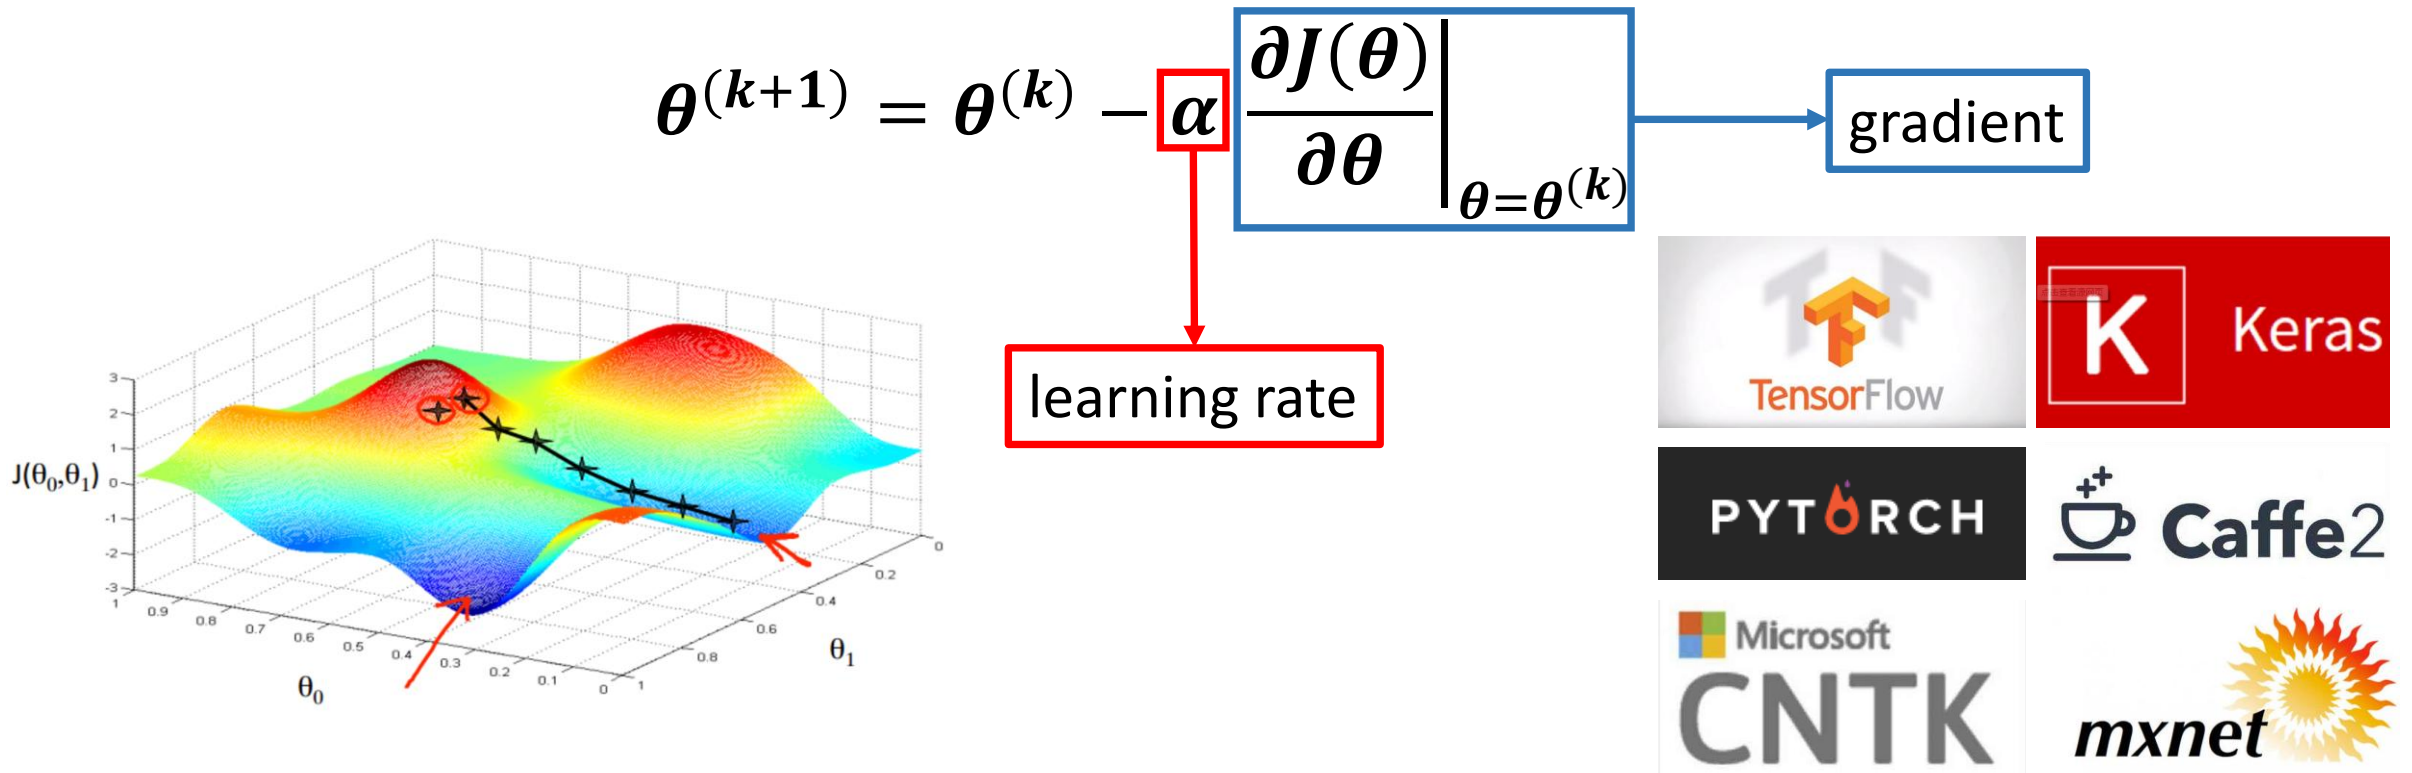
- Cross-entropy loss function:

$$J(\theta) = - \sum_{i \in D} \sum_{j=1}^k 1\{y_i = k\} \ln(f_k(x_i; \theta))$$

Neural Network: Feedforward propagation

- It's very difficult to find the **minimum** by traditional **gradient-based** method.

So we utilize **gradient descent algorithm**:



Neural Network: Feedforward propagation

gradient descent algorithm

$$j(x_i; \theta) = \frac{1}{2} \|f(x_i; \theta) - y_i\|^2$$

Batch Gradient Descent

$$\theta^{(k+1)} = \theta^{(k)} - \alpha \sum_{i=1}^N \frac{\partial j(x_i; \theta)}{\partial \theta} \Big|_{\theta=\theta^{(k)}}$$


Stochastic Gradient Descent

$$\theta^{(k+1)} = \theta^{(k)} - \alpha \frac{\partial j(\mathbf{x}_i; \theta)}{\partial \theta} \Big|_{\theta=\theta^{(k)}}$$

Mini-batch Gradient Descent $\mathbf{D} = [\mathbf{D}_1, \mathbf{D}_2, \dots, \mathbf{D}_M]$

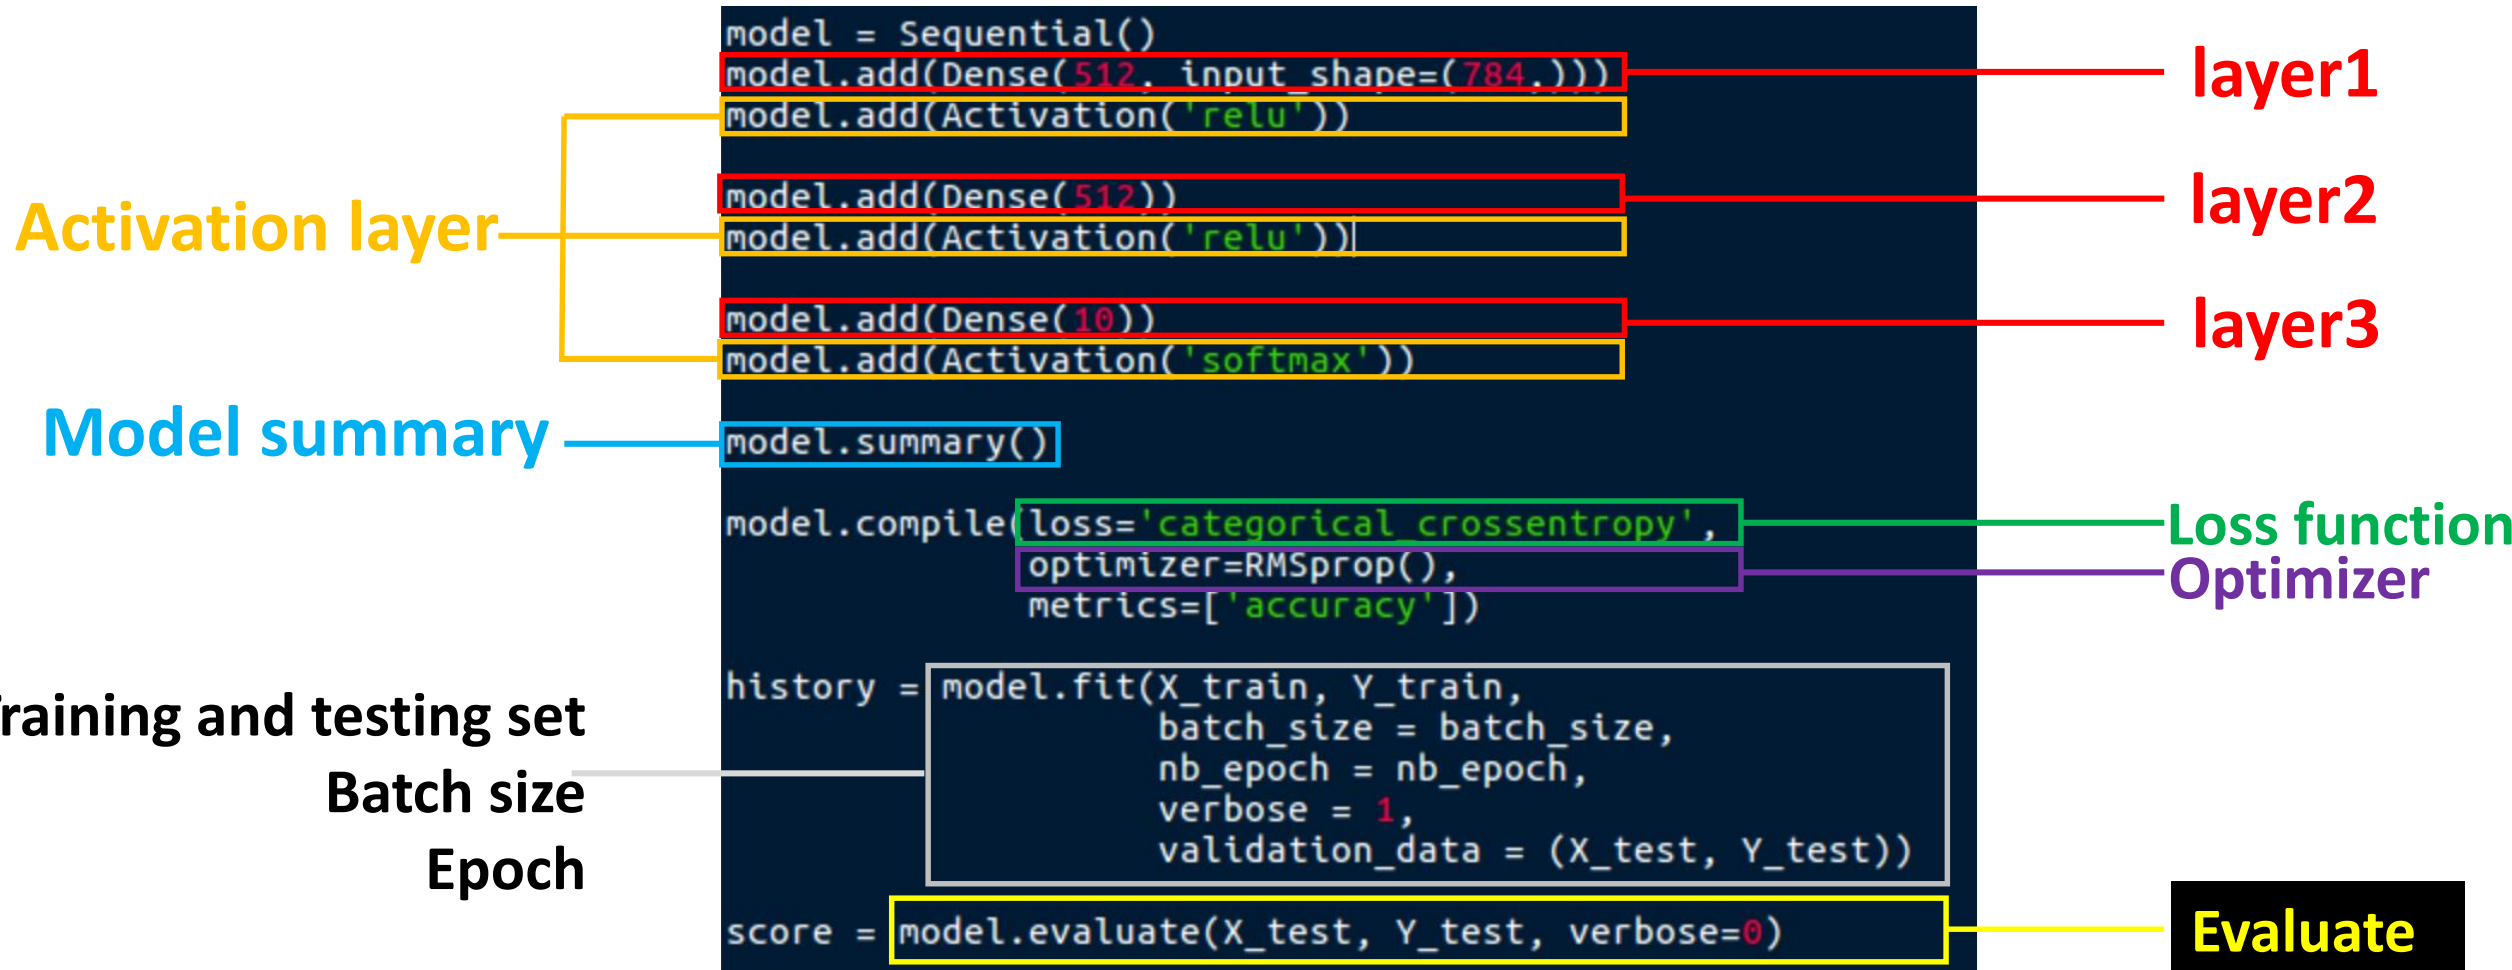
$$\theta^{(k+1)} = \theta^{(k)} - \alpha \sum_{i \in \mathbf{D}_i} \frac{\partial j(x_i; \theta)}{\partial \theta} \Big|_{\theta=\theta^{(k)}}$$

Neural Network: Summary



- Step1 : **Construct** a neural network, confirm the number of neurons in each layer.
- Step2: Confirm the **loss function** : MSE, Cross-entropy.
- Step3: Confirm the **optimization method**: batch size and epochs.
- Step4: **Random initialize** the parameters (weights and bias).
- Step5: Divide the dataset into two part, **train: test**=4:1 or 9:1.
- Step6: **Training** model.
- Step7: **Evaluate** model, by loss or accuracy.

Neural Network: Experiment



Neural Network: Result

20 neurons

vs

512neurons

epoch1	loss: 0.5601 - acc: 0.8527 - val_loss: 0.3027 - val_acc: 0.9148	loss: 0.2734 - acc: 0.9208 - val_loss: 0.1291 - val_acc: 0.9622
epoch2	loss: 0.2874 - acc: 0.9183 - val_loss: 0.2617 - val_acc: 0.9256	loss: 0.1179 - acc: 0.9653 - val_loss: 0.1007 - val_acc: 0.9694
epoch3	loss: 0.2540 - acc: 0.9273 - val_loss: 0.2380 - val_acc: 0.9300	loss: 0.0819 - acc: 0.9749 - val_loss: 0.0764 - val_acc: 0.9769
epoch4	loss: 0.2300 - acc: 0.9343 - val_loss: 0.2193 - val_acc: 0.9368	loss: 0.0635 - acc: 0.9810 - val_loss: 0.0757 - val_acc: 0.9763
epoch5	loss: 0.2108 - acc: 0.9395 - val_loss: 0.2152 - val_acc: 0.9363	loss: 0.0507 - acc: 0.9843 - val_loss: 0.0630 - val_acc: 0.9809
epoch6	loss: 0.1963 - acc: 0.9439 - val_loss: 0.1940 - val_acc: 0.9414	loss: 0.0426 - acc: 0.9871 - val_loss: 0.0636 - val_acc: 0.9802
epoch7	loss: 0.1853 - acc: 0.9466 - val_loss: 0.1952 - val_acc: 0.9436	loss: 0.0355 - acc: 0.9891 - val_loss: 0.0633 - val_acc: 0.9816
epoch8	loss: 0.1764 - acc: 0.9488 - val_loss: 0.1835 - val_acc: 0.9460	loss: 0.0325 - acc: 0.9904 - val_loss: 0.0616 - val_acc: 0.9830
epoch9	loss: 0.1691 - acc: 0.9511 - val_loss: 0.1856 - val_acc: 0.9452	loss: 0.0272 - acc: 0.9918 - val_loss: 0.0625 - val_acc: 0.9819
epoch10	loss: 0.1629 - acc: 0.9525 - val_loss: 0.1747 - val_acc: 0.9494	loss: 0.0233 - acc: 0.9925 - val_loss: 0.0640 - val_acc: 0.9816
	val_loss: 0.1747 - val_acc: 0.9494	val_loss: 0.0640 - val_acc: 0.9816

Content

~~• Neural Network~~

• Convolution Neural Network

• Deep Learning for Computer Vision

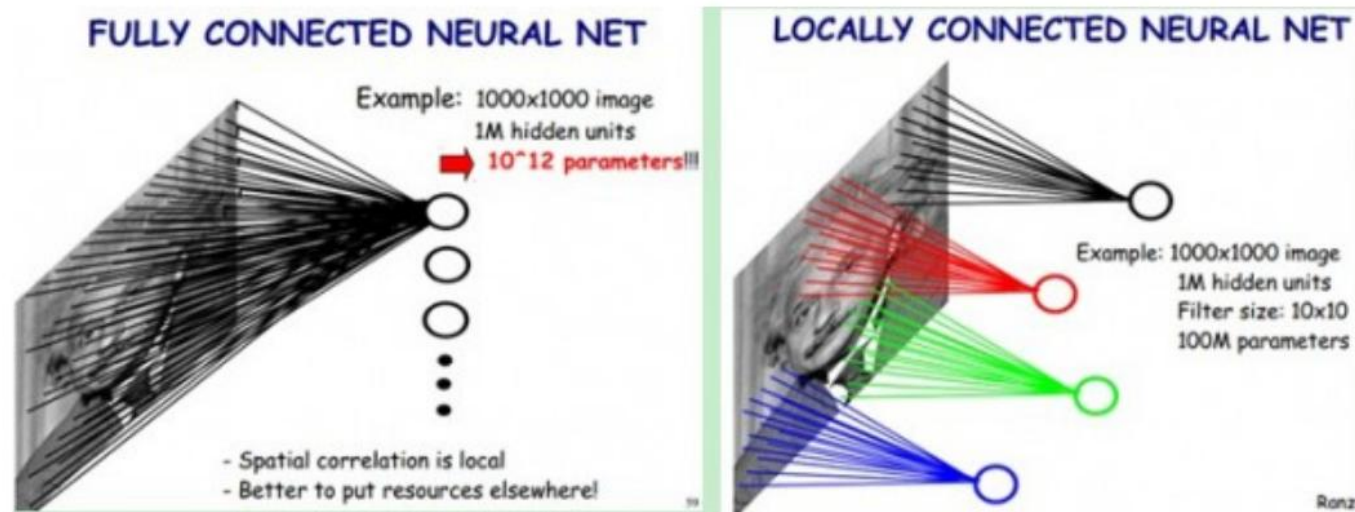
Motivation

Convolutional layer

Max-pooling layer

Classical CNN

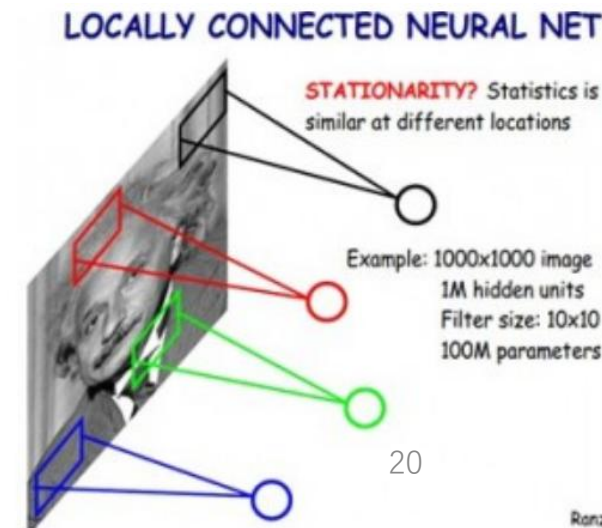
Convolutional Neural Network: Motivations



- There are two approach to reduce the number of parameters:
- **Receptive field.**
- **Sharing weights.**




Convolutional Neural Network: Convolution layer

- **Fully connected network:**
- With larger images (e.g., **96x96 images**), there are about **10^4 input units**, and assuming you want to learn **100 features**, you would have on the order of **10^6** parameters to learn.
- **Locally connected network:**
- Each hidden unit will connect **to only a small contiguous region** of pixels in the input.



Convolutional Neural Network: Convolution layer

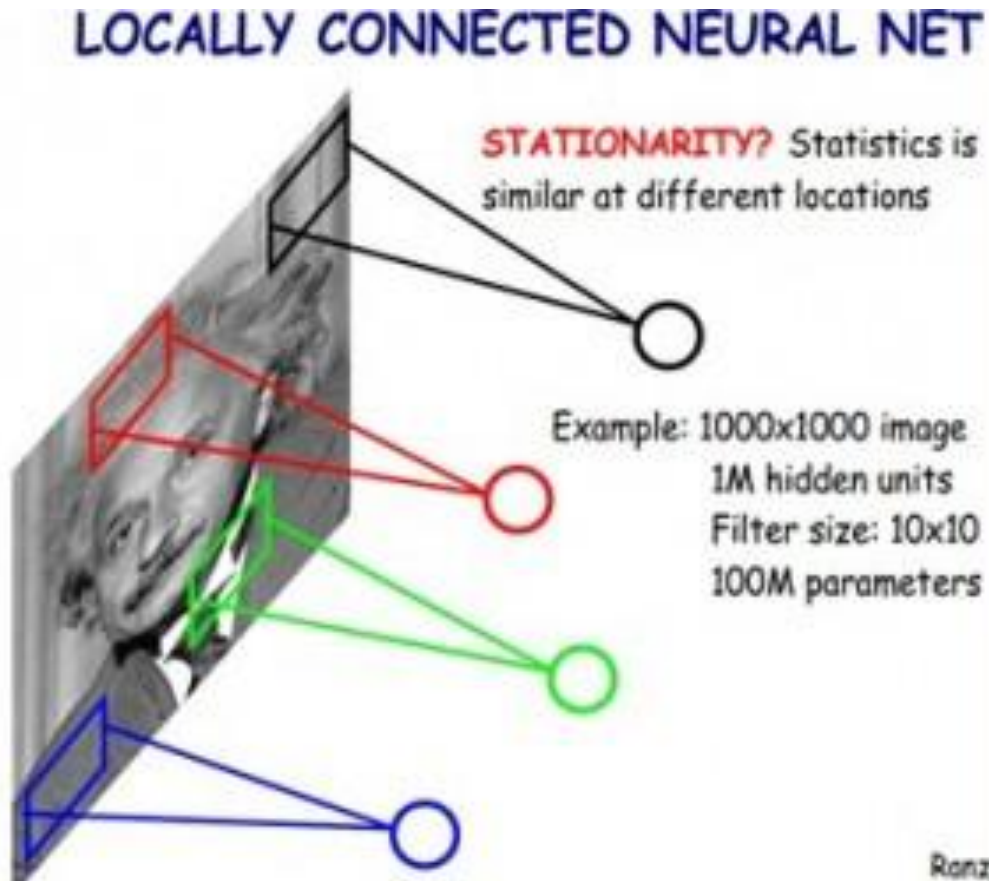
- Natural images have the property of being **stationary**



- The statistics of one part of the image are the **same as any other part**.
- This suggests that the features that we learn at one part of the image can also be applied to other parts of the image, and **we can use the same features at all locations**.


Convolutional Neural Network: Convolution layer

Locally connected



without sharing weights

vs



with sharing weights

Convolutional Neural Network: Convolution layer

Convolution for single-channel:

$$f(x, y) \circ w = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) \cdot f(x-s, y-t)$$




rotate the kernel 180° & correlation

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

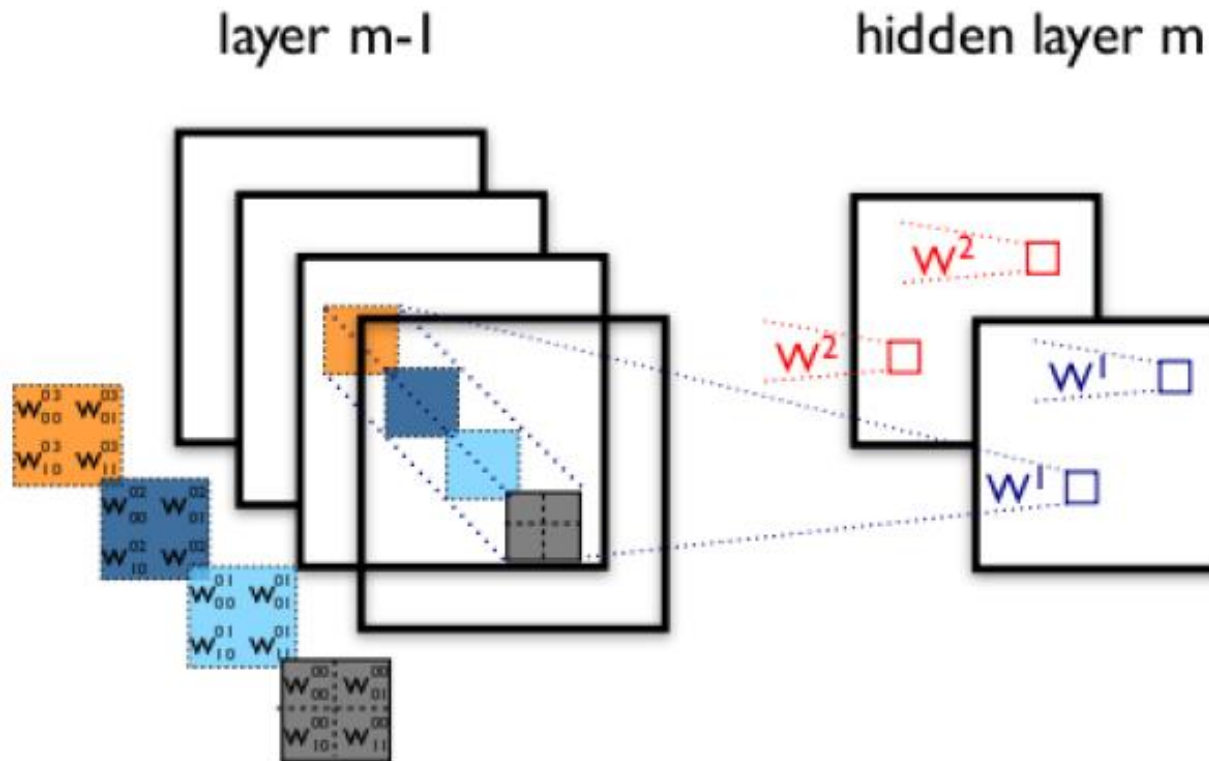
Convolved
Feature



Convolutional Neural Network: Convolution layer

Convolution for multi-channel:



Supposing that there are n_{m-1} channels in layer m-1 and n_m channels in layer m and there are kernels with size $S \times t$, the number of parameters of layer m-1 and layer m is $n_{m-1} \times n_m \times S \times t + n_m$.



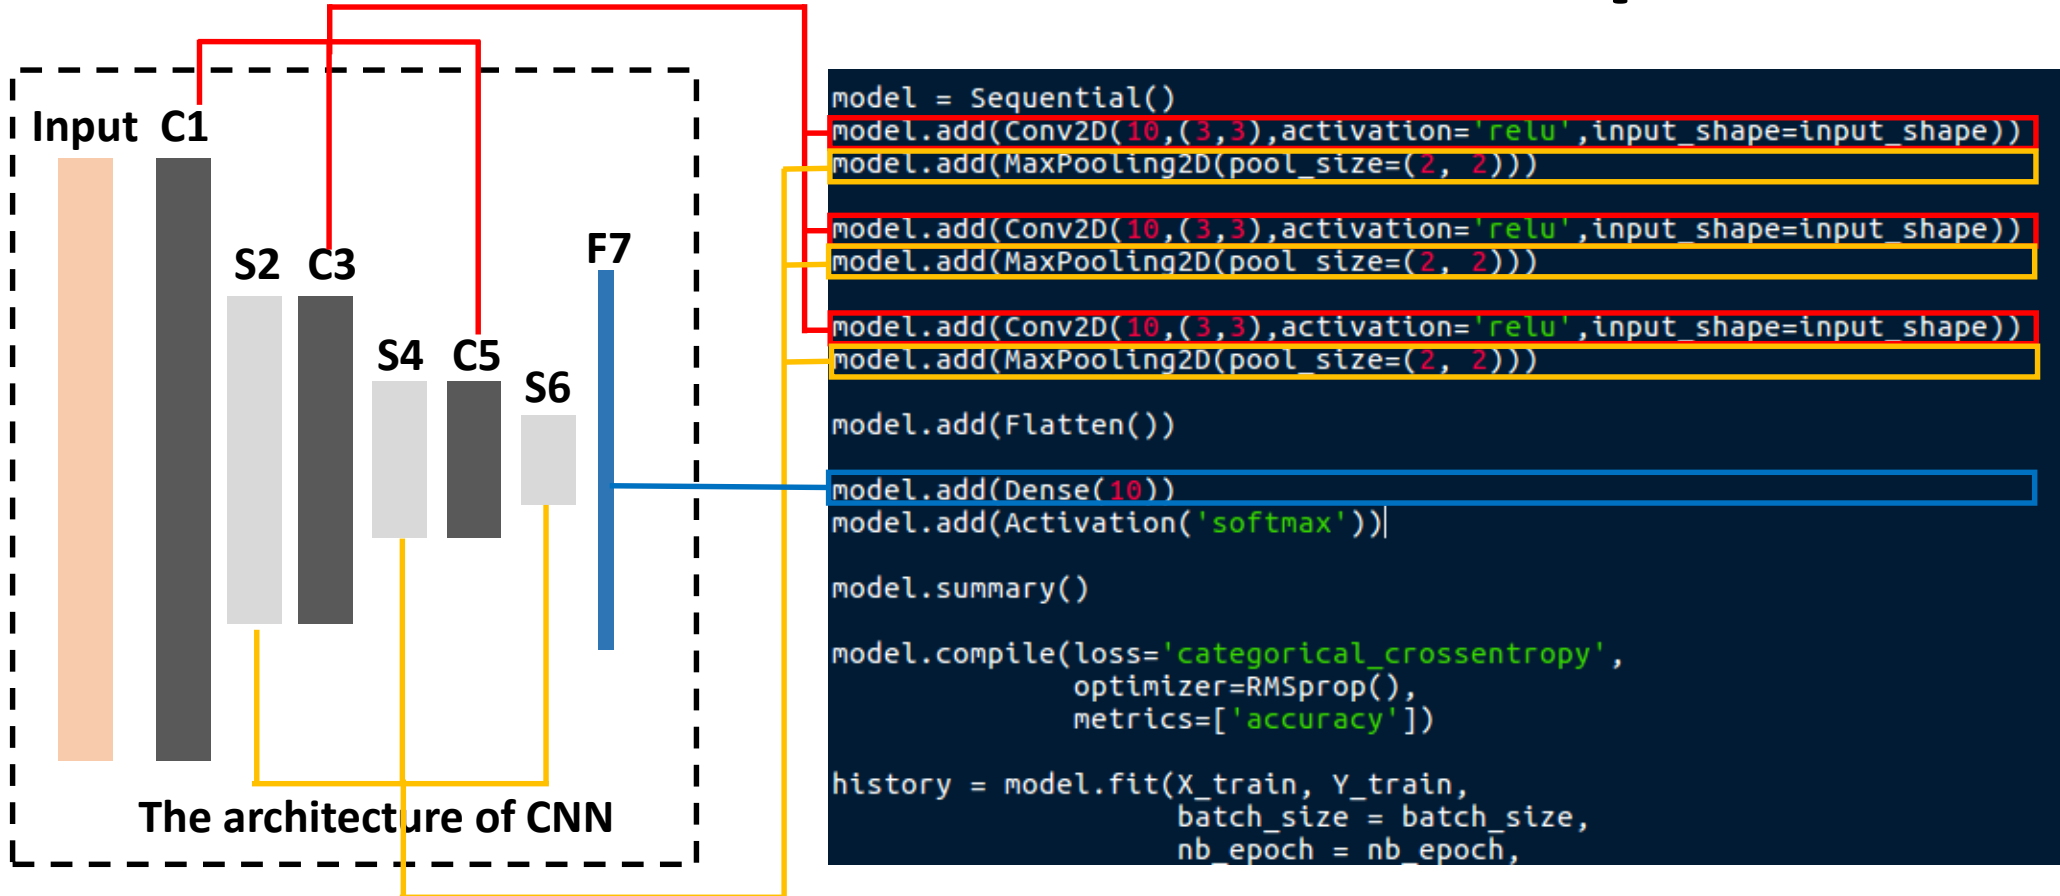
The number of parameters of layer m-1 and layer m is $4 \times 2 \times 2 \times 2 + 2 = 34$.

Convolutional Neural Network: Max-pooling Layer

- For a input gray image with $96*96$, if we use $400\ 8*8$ kernels to obtain feature maps, we will get 400 feature maps with size $(96-8+1)*(96-8+1)$.
- It means we have $400*(96-8+1)*(96-8+1)=3168400$ features!!
- Max-pooling
- Mean-pooling



Convolutional Neural Network: Experiment



Convolutional Neural Network: Experiment

20 neurons

vs

CNN

epoch1	loss: 0.5601 - acc: 0.8527 - val_loss: 0.3027 - val_acc: 0.9148	loss: 1.0958 - acc: 0.6510 - val_loss: 0.4778 - val_acc: 0.8579
epoch2	loss: 0.2874 - acc: 0.9183 - val_loss: 0.2617 - val_acc: 0.9256	loss: 0.3964 - acc: 0.8837 - val_loss: 0.2943 - val_acc: 0.9148
epoch3	loss: 0.2540 - acc: 0.9273 - val_loss: 0.2380 - val_acc: 0.9300	loss: 0.2926 - acc: 0.9128 - val_loss: 0.2399 - val_acc: 0.9277
epoch4	loss: 0.2300 - acc: 0.9343 - val_loss: 0.2193 - val_acc: 0.9368	loss: 0.2453 - acc: 0.9257 - val_loss: 0.2042 - val_acc: 0.9392
epoch5	loss: 0.2108 - acc: 0.9395 - val_loss: 0.2152 - val_acc: 0.9363	loss: 0.2138 - acc: 0.9355 - val_loss: 0.2097 - val_acc: 0.9372
epoch6	loss: 0.1963 - acc: 0.9439 - val_loss: 0.1940 - val_acc: 0.9414	loss: 0.1895 - acc: 0.9424 - val_loss: 0.1813 - val_acc: 0.9436
epoch7	loss: 0.1853 - acc: 0.9466 - val_loss: 0.1952 - val_acc: 0.9436	loss: 0.1724 - acc: 0.9477 - val_loss: 0.1529 - val_acc: 0.9545
epoch8	loss: 0.1764 - acc: 0.9488 - val_loss: 0.1835 - val_acc: 0.9460	loss: 0.1569 - acc: 0.9523 - val_loss: 0.1454 - val_acc: 0.9564
epoch9	loss: 0.1691 - acc: 0.9511 - val_loss: 0.1856 - val_acc: 0.9452	loss: 0.1440 - acc: 0.9558 - val_loss: 0.1523 - val_acc: 0.9540
epoch10	loss: 0.1629 - acc: 0.9525 - val_loss: 0.1747 - val_acc: 0.9494	loss: 0.1354 - acc: 0.9582 - val_loss: 0.1236 - val_acc: 0.9620

val_loss: 0.1747 - val_acc: 0.9494

val_loss: 0.1236 - val_acc: 0.9620

$28*28*20+20 = 15700$ parameters

vs

2030 parameters

Classical CNN model

- **Lenet——1998**
- **Alexnet——2012**
- **VGG-net——2014**


Classical CNN model——Lenet

- 7 layers:
- Input layer: 1@32*32
- layer1: 6@5*5 kernel
- layer2: 2*2 max-pooling
- layer3: 16 feature map

feature map1: 6@28*28

feature map2: 6@14*14

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	X				X	X	X			X	X	X	X		X	X
1	X	X				X	X	X			X	X	X	X		X
2	X	X	X				X	X	X			X		X	X	X
3		X	X	X			X	X	X	X			X		X	X
4			X	X	X			X	X	X	X		X	X		X
5				X	X	X			X	X	X	X		X	X	X



Classical CNN: Lenet


- **layer4: max-pooling**
- **layer5: fully connected-120**
- **layer6: fully connected-84**
- **layer7: Gaussian connections**

feature map4: 16@5*5

feature map: 120

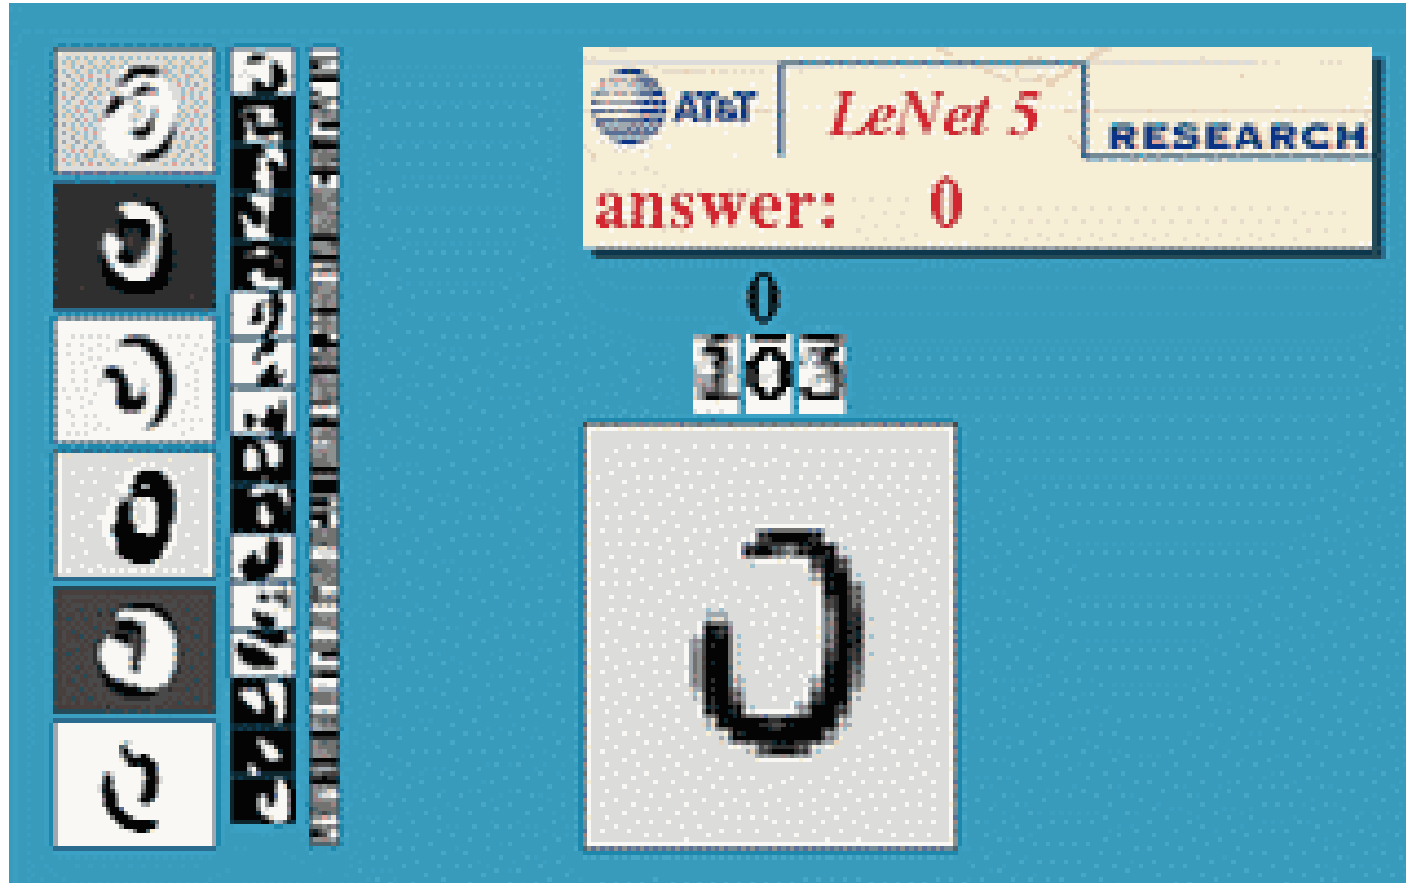
featuremap:84

feature map:10




Classical CNN: Lenet

- Demo from <http://yann.lecun.com/exdb/lenet/multiples.html>



Classical CNN model——Alexnet

- The champion of LSVRC2010, **top 5 error:15.3%, top 1 error:37.5%.**
- 1.2 million high resolution image, 1000 classes.
- **60 million parameters , 650,000 neurons, 8 learning layers.**



Classical CNN: Alexnet

11 layers:

Input layer: 3@224*224

layer1:96@11*11 kernels, stride4

layer2:max-pooling 3*3, stride 2

layer3:256@ 5*5 kernels

layer4:max-pooling 3*3, stride 2

layer5:384@ 5*5 kernels


feature map1:96@55*55

feature map2:96@27*27

feature map3: 256@27*27

feature map4:256@13*13

feature map5: 384@13*13



Classical CNN: Alexnet

layer6:384@ 5*5 kernels

layer7:256@ 5*5 kernels

layer8:max-pooling 3*3, stride 2

layer9:fc7-4096

layer10:fc7-4096

fc11+softmax: output 1000


feature map6: 384@13*13

feature map7: 256@13*13

feature map8:256@6*6

feature map9: 4096

feature map10: 4096



Classical CNN model——VGGnet

- The champion of LSVRC2014, **top 5 error:6.8%, top 1 error:23.7%.**

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

conv3-64

Why did they choose kernel with size **3*3** rather than 7*7 or 11*11?

Why did they choose kernel with size **3*3** rather than 7*7 or 11*11?

- **Reducing paramter.**

- A stack of two 3×3 conv. layers (without spatial pooling in between) has an effective receptive field of 5×5 .
- Three such layers have a 7×7 effective receptive field.

- **Increasing the nonlinearity of the decision function.**

- With 3 Relu nonlinear function.

Content

~~• Neural Network~~

~~• Convolution Neural Network~~

• Deep Learning for Computer Vision



- AutoEncoder
- Denoising AutoEncoder

Deep Learning for Computer Vision: AutoEncoder


- So far, we have described the application of neural networks to **supervised learning**, in which we have **labeled** training examples.
- Now suppose we have only a set of unlabeled training data $\{x^{(1)}, x^{(2)}, x^{(3)}, \dots\}$. An autoencoder neural network is an **unsupervised learning algorithm** that applies backpropagation, **setting the target values to be equal to the inputs.**



$$y^{(i)} = x^{(i)}$$

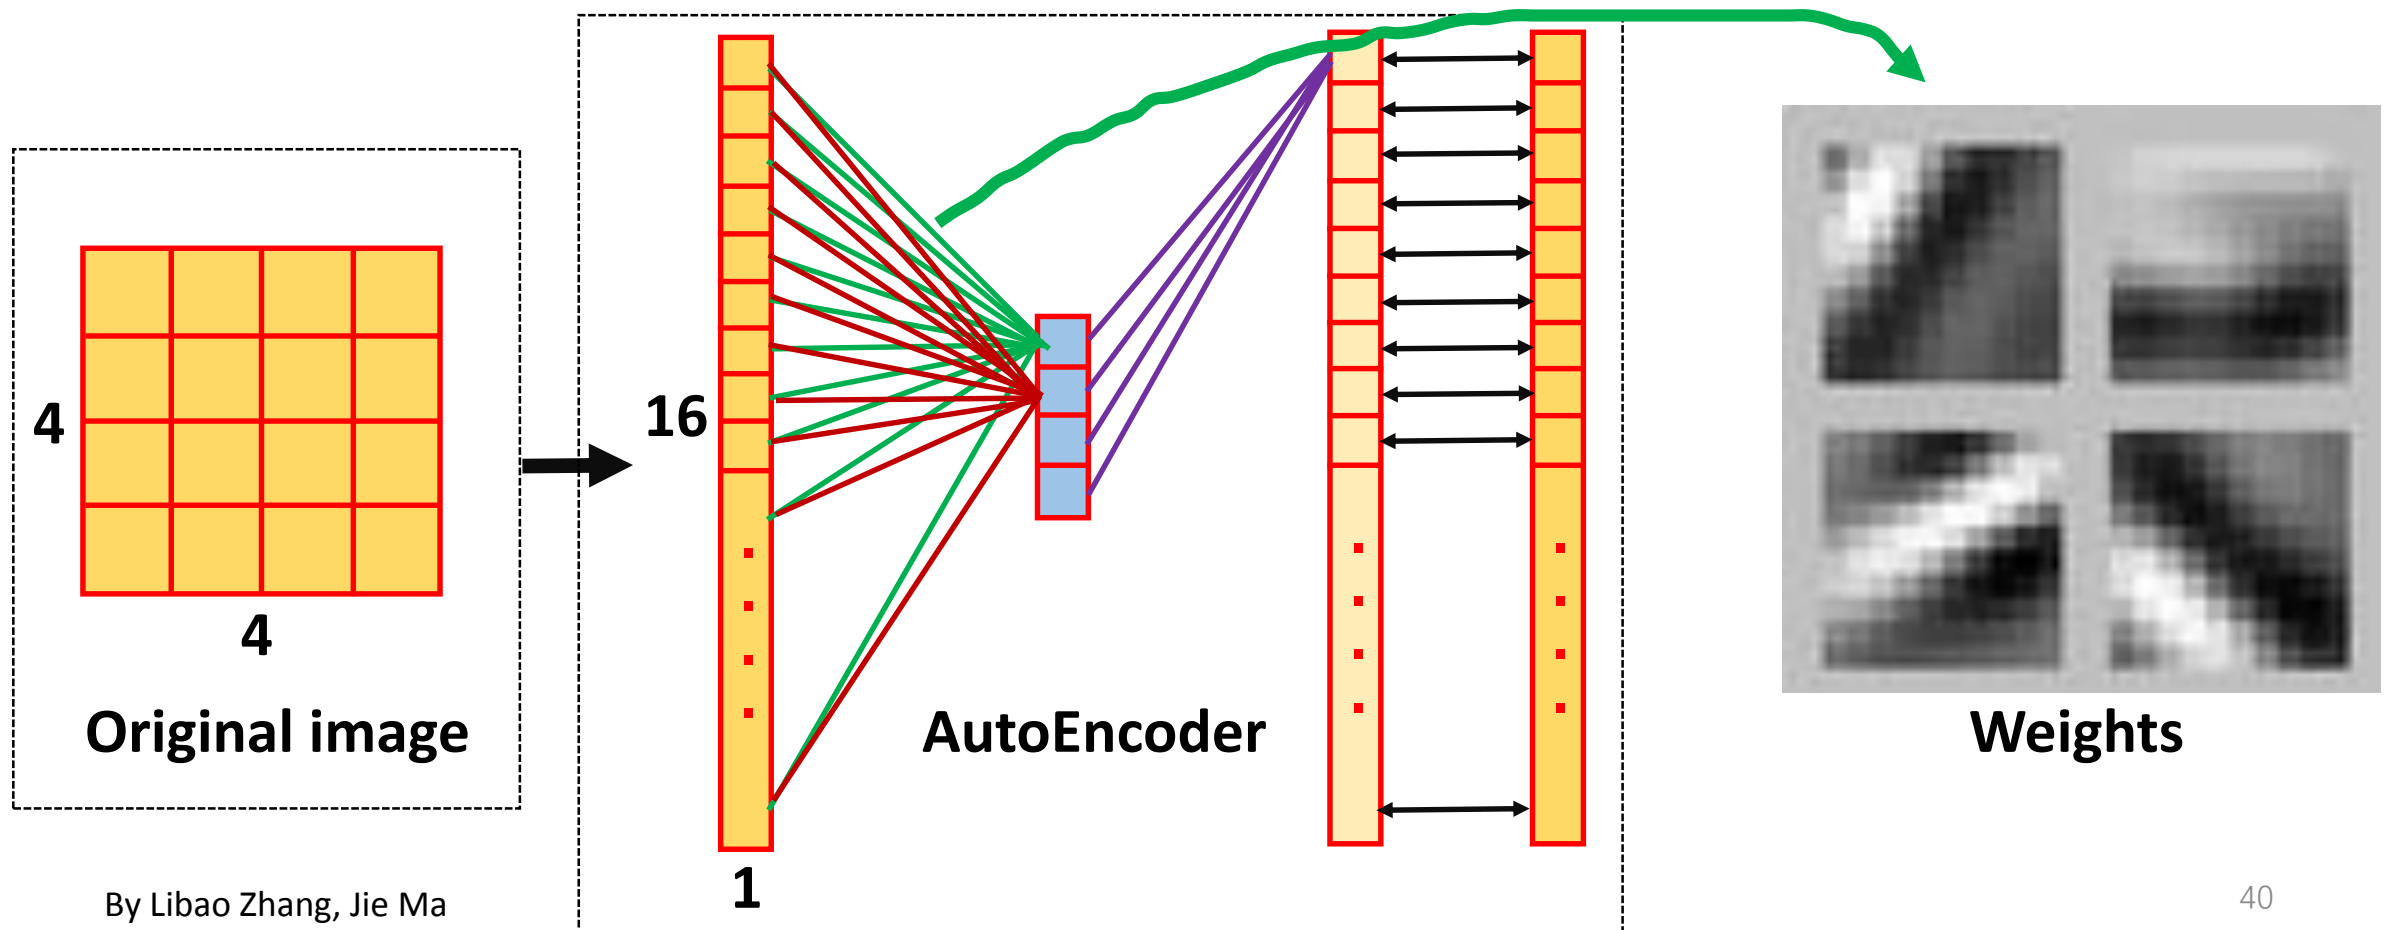
Deep Learning for Computer Vision : AutoEncoder

- The AutoEncoder tries to learn a function $f(x, \theta) = x$
- Suppose the inputs x are the pixel intensity values from a 10×10 image. so $n_1 = 100$, and there are $n_2 = 50$ hidden units in layer2.
- Note that we also have $y \in \mathbb{R}^{100}$. Since there are only **50 hidden units**, the network is forced to learn a **compressed representation** of the input.




Deep Learning for Computer Vision : AutoEncoder

- Visualizing a trained Autoencoder:
- We will visualize the function computed by **hidden unit i** ——**which depends on the parameters weight $W^{(1)}_{ij}$**



Deep Learning for Computer Vision : AutoEncoder

- **Stack** AutoEncoder



Deep Learning for Computer Vision : AutoEncoder

- **Stack** AutoEncoder: Experiment

```
model = Sequential()
model.add(Conv2D(10,(3,3),activation='relu',border_mode='same',input_shape=input_shape)) #10@28*28
model.add(MaxPooling2D(pool_size=(2, 2))) #10@14*14

model.add(Conv2D(20,(3,3),activation='relu',border_mode='same',input_shape=input_shape)) #10@14*14
model.add(MaxPooling2D(pool_size=(2, 2))) #10@7*7

model.add(Conv2D(30,(3,3),activation='relu',border_mode='same',input_shape=input_shape)) #10@7*7

model.add(UpSampling2D(size=(2, 2))) #10@14*14
model.add(Conv2D(20,(3,3),activation='relu',border_mode='same',input_shape=input_shape)) #10@14*14

model.add(UpSampling2D(size=(2, 2))) #10@28*28
model.add(Conv2D(10,(3,3),activation='relu',border_mode='same',input_shape=input_shape)) #10*28*28

model.add(Conv2D(1,(3,3),activation='relu',border_mode='same',input_shape=input_shape))

model.compile(loss='mse',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(X_train, X_train,
                   batch_size = batch_size,
                   nb_epoch = nb_epoch,
                   verbose = 1,
                   validation_data = (X_test, X_test))
```


Encoder

Decoder

Deep Learning for Computer Vision : AutoEncoder

- **Stack** AutoEncoder: Experiment

Original image



decode image



Deep Learning for Computer Vision: Denoising AutoEncoder


- Question:

How to learn **more robust** features?


- In order to obtain more **robust feature**, we use **damaged input** to **reconstruct** the image.



Deep Learning for Computer Vision : Denosing AutoEncoder



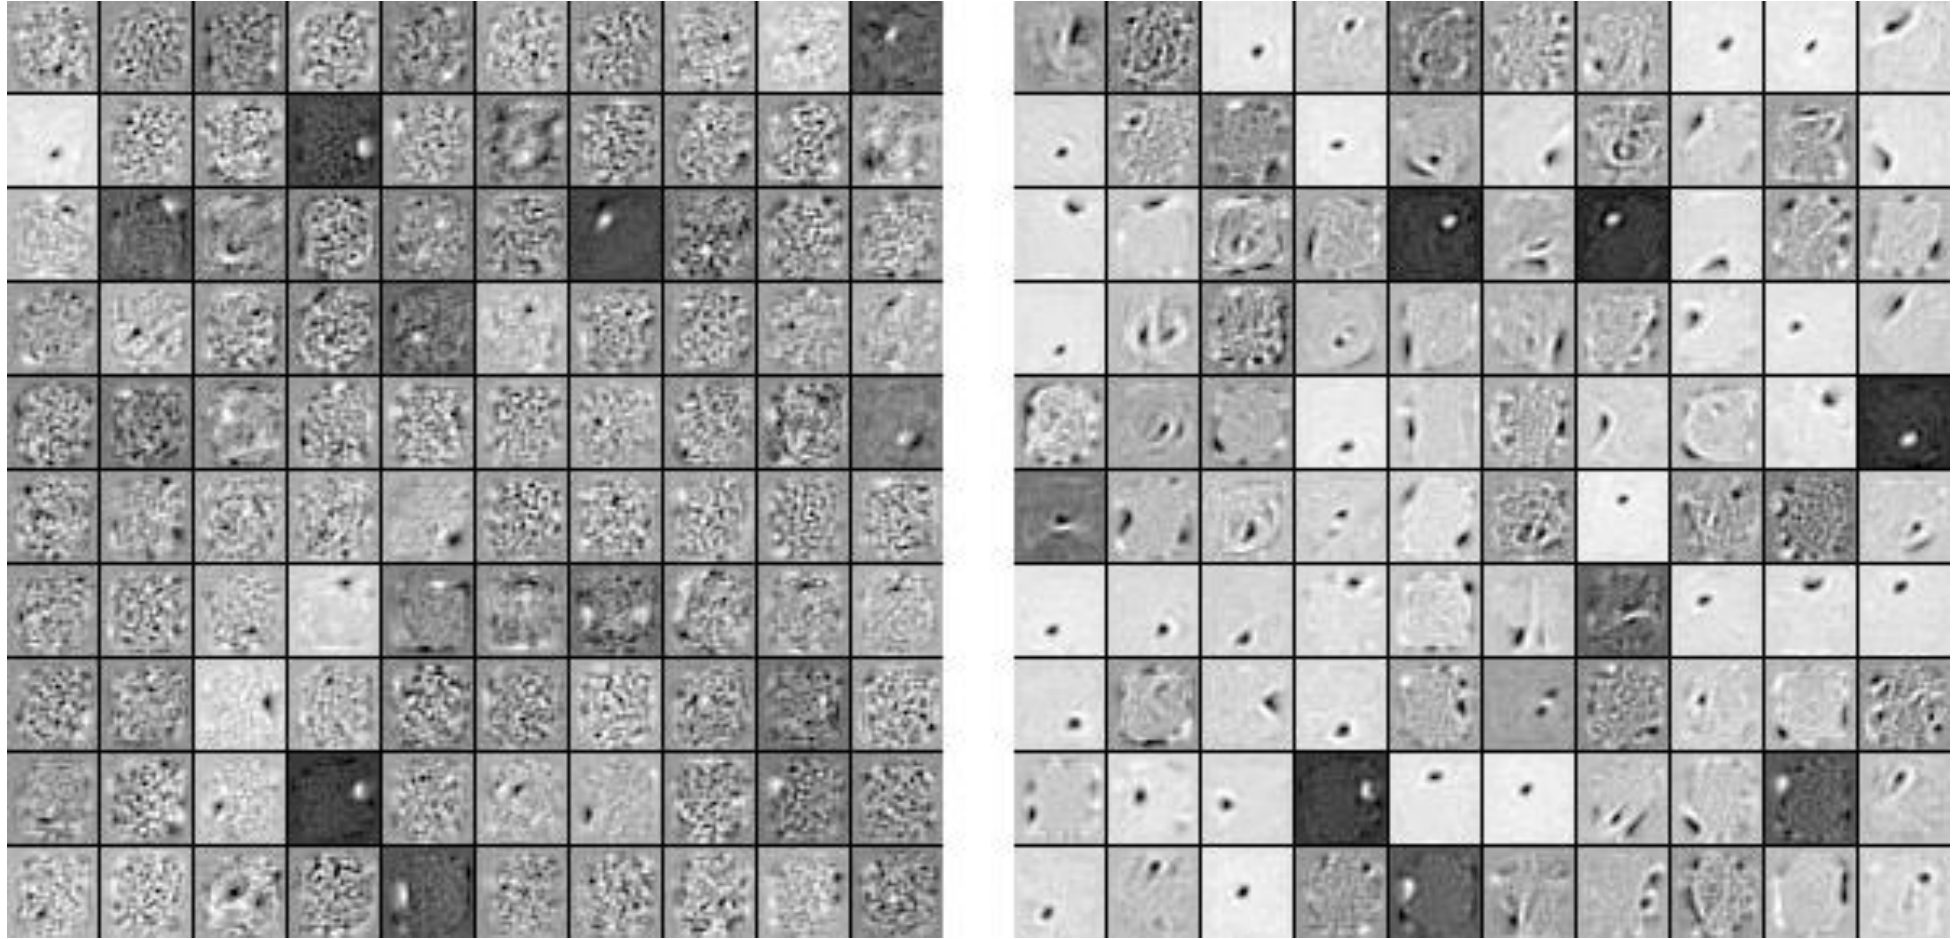
AutoEncoder



**Denosing
AutoEncoder**

Deep Learning for Computer Vision : Denosing AutoEncoder

- Comparision between **AutoEncoder** and **Denosing AutoEncoder**.






Summary

- Neural Network
 - Discussion from XOR
 - Forward Propagation
 - Back Propagation

- Convolution Neural Network
 - Motivation
 - Convolutional layer
 - Max-pooling layer
 - Classical CNN
 - Lenet
 - Alexnet
 - VGG

- Deep Learning for Computer Vision
 - AutoEncoder
 - Stack AutoEncoder
 - Denoise AutoEncoder



Public email address

deeplearning_class@163.com